

lesson5_naivebayes

April 27, 2018

```
In [1]: # Run this every time you open the spreadsheet
        %load_ext autoreload
        %autoreload 2
        from collections import Counter
        import lib
```

1 Load and inspect the data

```
In [2]: # Load the data.
        # This function returns tweets and test_tweets, both lists of tweets
        tweets, test_tweets = lib.read_data()
```

2 Learn a Naive Bayes classifier

To construct our Naive Bayes classifier, we first need to calculate two things:

2.0.1 Prior probabilities of categories

We need to calculate $P(C_i)$ for each category $C_i \in \{\text{Energy, Food, Medical, Water, None}\}$.

We estimate $P(C_i)$ by $\frac{\# \text{ tweets about } C_i}{\# \text{ tweets}}$

2.0.2 Conditional probabilities of tokens

For each token (i.e. word) x_j and each category C_i , we need to calculate $P(x_j|C_i)$.

We estimate $P(x_j|C_i) = \frac{P(x_j \text{ and } C_i)}{P(C_i)}$ by $\frac{\# \text{ tweets about } C_i \text{ containing } x_j}{\# \text{ tweets about } C_i}$

```
In [3]: # Exercise 1, step-by-step version (challenge version is below).
```

```
# The function below has two arguments: a list of tweets, and a category c
# which is a string equal to one of "Energy", "Food", "Medical", "Water", "None".
# The function should calculate the two things described above.
# Fill in the blanks.
```

```
def calc_probs(tweets, c):
    """
```

Input:

tweets: a list of tweets

c: a string representing a category; one of "Energy", "Food", "Medical", "Water"

Returns:

prob_c: the prior probability of category c

token_probs: a Counter mapping each token to $P(\text{token}/\text{category } c)$

"""

Step 1: Calculate the total number of tweets

`num_tweets = len(tweets)`

Step 2: Calculate the number of tweets that are about category c.

Save the answer to a variable called num_tweets_about_c.

Remember c is a string, and you can get the category of a tweet via tweet.category

`num_tweets_about_c = sum(map(lambda tweet: tweet.category == c, tweets))`

Step 3: Calculate the probability of category c using the answers from Steps 1 and 2.

Hint: be careful when you divide two integers!

`prob_c = float(num_tweets_about_c)/num_tweets`

Step 4: Create an empty Counter called token_counts.

(We will use it to map each token to the number of category-c tweets containing it)

`token_counts = Counter()`

Step 5 (tricky): Use a for-loop to iterate over the list of tweets.

Use an if-statement to check whether the tweet is in category c.

If it is, iterate over the tokens of the tweet (which you can access via tweet.tokens)

For each token, increment its count in token_counts.

`for tweet in tweets:`

`if tweet.category == c:`

`for token in tweet.tokens:`

`token_counts[token] += 1`

Step 6: Create an empty Counter called token_probs.

(We will use it to map each token to $P(\text{token} | \text{category } c)$,

i.e. the fraction of all category-c tweets that contain the token)

`token_probs = Counter()`

Step 7: Now fill token_probs.

For each token->count in token_counts, you want to add token->fraction to token_probs[token]

Use a for-loop over token_counts.

Remember that when you iterate over a dictionary/Counter, you access the keys.

You'll need to use the variable num_tweets_about_c.

Be careful when you divide integers!

`for token in token_counts:`

```

token_probs[token] = token_counts[token] / num_tweets_about_c

print("Class %s has prior probability %.2f" % (c, prob_c))
return prob_c, token_probs

prob_food, token_probs_food = calc_probs(tweets, "Food")
prob_water, token_probs_water = calc_probs(tweets, "Water")
prob_energy, token_probs_energy = calc_probs(tweets, "Energy")
prob_medical, token_probs_medical = calc_probs(tweets, "Medical")
prob_none, token_probs_none = calc_probs(tweets, "None")

Class Food has prior probability 0.47
Class Water has prior probability 0.09
Class Energy has prior probability 0.12
Class Medical has prior probability 0.04
Class None has prior probability 0.28

```

2.0.3 See what your model has learnt

In [4]: *# For each category c, print out the tokens that maximize P(c/token)*

```

token_probs = {'Food': token_probs_food, 'Water': token_probs_water, 'Energy': token_p
               'Medical': token_probs_medical, 'None': token_probs_none}
prior_probs = {'Food': prob_food, 'Water': prob_water, 'Energy': prob_energy, 'Medical
               'None': prob_none}
lib.most_discriminative(tweets, token_probs, prior_probs)

```

MOST DISCRIMINATIVE TOKENS:

TOKEN	P(Energy token)
powers	0.8029
dark	0.8029
generator	0.7654
batteries	0.7559
class	0.7534
sandysucks	0.7534
flashlights	0.7345
masks	0.7334
11/3	0.6736
cleaner	0.6707

TOKEN	P(Food token)
canned	0.9784
non-perishable	0.9767
serve	0.9663
perishable	0.9562

cook	0.9511
soup	0.9489
sandwiches	0.9489
thanksgiving	0.9441
rice	0.9441
meal	0.9383

TOKEN	P(Medical token)
meds	0.8229
aid	0.8008
ointment	0.7360
prescription	0.7360
ups	0.7360
medicine	0.7360
medications	0.7360
4t-5t	0.7360
kits	0.6596
pull	0.6596

TOKEN	P(None token)
..	0.9531
everyone	0.8955
last	0.8809
feel	0.8809
im	0.8618
irene	0.8604
...	0.8601
thing	0.8314
wow	0.8314
tropical	0.8314

TOKEN	P(Water token)
bottled	0.9059
gallon	0.8307
jugs	0.7970
water	0.7873
gallons	0.7266
pallets	0.6625
spring	0.6625
flood	0.6625
liter	0.6625
parks	0.6625

3 Build a Naive Bayes classifier

Now we've calculated $P(C_i)$ and $P(x_j|C_i)$, we can classify any tweet!

Given a tweet which is a set of tokens $\{x_1, \dots, x_n\}$, the posterior probability of each category C_i is

$$P(C_i|x_1, \dots, x_n) \propto P(C_i) \times P(x_1|C_i) \times P(x_2|C_i) \dots \times P(x_n|C_i)$$

We just need to calculate this for each category then determine which is largest.

In [5]: # Exercise 2.

```
# Complete this function that calculates the posterior probability of P(c/tweet).
```

```
def get_posterior_prob(tweet, prob_c, token_probs):  
    """Calculate the posterior P(c/tweet).  
    (Actually, calculate something proportional to it).  
  
    Inputs:  
        tweet: a tweet  
        prob_c: the prior probability of category c  
        token_probs: a Counter mapping each token P(token/c)  
    Return:  
        The posterior P(c/tweet).  
    """
```

```
##### YOUR CODE STARTS HERE #####
```

```
# Hint: first set posterior to prob_c, then use a for-loop over tweet.tokenSet  
# to repeatedly multiply posterior by P(token/c)
```

```
posterior = prob_c  
for token in tweet.tokenSet:  
    if token_probs[token] == 0:  
        posterior *= 0.001  
    else:  
        posterior *= token_probs[token]
```

```
##### YOUR CODE ENDS HERE #####
```

```
return posterior
```

```
# Now you've written the function, look at the output for P(Energy/"No power in Riverdale")  
# What's gone wrong?  
# Try editing your function above to print out each token and token_probs[token].  
# Can you see what went wrong? How might you fix it?
```

```
riverdale_tweet = lib.Tweet("No power in Riverdale", "Energy", "need")
```

```
print("P(Energy|'No power in Riverdale') = ", get_posterior_prob(riverdale_tweet, prob_
P(Energy|'No power in Riverdale') = 2.806001890359169e-06
```

```
In [6]: # This cell defines the classification function, that takes a tweet
# and decides which category is most likely using the posteriors you just calculated.
```

```
# OPTIONAL EXERCISE (come back to it once you've reached the end of the notebook).
# Rewrite this function to be less repetitive i.e. don't repeat things 5 times.
# There are several possible solutions; you might want to use lists or dictionaries.
# You might also want to rewrite the earlier code that computed prob_food, token_probs
```

```
def classify_nb(tweet):
    """Classifies a tweet. Calculates the posterior P(c|tweet) for each category c,
    and returns the category with largest posterior.
    Input:
        tweet
    Output:
        string equal to most-likely category for this tweet
    """
    posterior_food_prob = get_posterior_prob(tweet, prob_food, token_probs_food)
    posterior_water_prob = get_posterior_prob(tweet, prob_water, token_probs_water)
    posterior_energy_prob = get_posterior_prob(tweet, prob_energy, token_probs_energy)
    posterior_medical_prob = get_posterior_prob(tweet, prob_medical, token_probs_medical)
    posterior_none_prob = get_posterior_prob(tweet, prob_none, token_probs_none)

    max_posterior = max([posterior_food_prob, posterior_water_prob,
                        posterior_energy_prob, posterior_medical_prob,
                        posterior_none_prob])
    if posterior_food_prob == max_posterior:
        return 'Food'
    elif posterior_water_prob == max_posterior:
        return 'Water'
    elif posterior_energy_prob == max_posterior:
        return 'Energy'
    elif posterior_medical_prob == max_posterior:
        return 'Medical'
    else:
        return 'None'
```

3.1 Evaluate the Naive Bayes classifier

```
In [7]: # Compare true labels and predicted labels in a table
```

```
predictions = [(tweet, classify_nb(tweet)) for tweet in test_tweets] # a list of (tweet, prediction)
lib.show_predictions(predictions)
```

<IPython.core.display.HTML object>

In [8]: *# Get average F1 score for the test set*

```
predictions = [(tweet, classify_nb(tweet)) for tweet in test_tweets] # maps each test  
lib.evaluate(predictions)
```

Energy

Precision: 50.0

Recall: 60.0

F1: 54.54545454545455

Food

Precision: 83.56164383561644

Recall: 94.57364341085271

F1: 88.72727272727272

Medical

Precision: 85.71428571428571

Recall: 46.15384615384615

F1: 60.0

None

Precision: 82.85714285714286

Recall: 73.41772151898734

F1: 77.85234899328859

Water

Precision: 80.0

Recall: 40.0

F1: 53.333333333333336

Average F1: 66.89168191986984

In [9]: *# Get average F1 score for the TRAINING set.*

Compare with average F1 for test set above. What's the reason for the difference?

```
trainset_predictions = [(tweet, classify_nb(tweet))  
                        for tweet in tweets] # maps each training tweet to its predic  
lib.evaluate(trainset_predictions)
```

Energy

Precision: 91.33333333333333

Recall: 99.27536231884058

F1: 95.13888888888887

Food

Precision: 96.6355140186916
Recall: 97.91666666666667
F1: 97.27187206020695

Medical

Precision: 97.77777777777777
Recall: 100.0
F1: 98.87640449438202

None

Precision: 97.98657718120805
Recall: 94.49838187702265
F1: 96.21087314662273

Water

Precision: 100.0
Recall: 91.08910891089108
F1: 95.33678756476684

Average F1: 96.56696523097348

```
In [10]: lib.show_confusion_matrix(predictions)
```

<IPython.core.display.HTML object>